

Installing CJE In Digital Ocean

The goal of this document is to describe how to install a Kubernetes cluster in **DigitalOcean** using **Rancher**. Unlike AWS, GCE, and Azure, DigitalOcean does not offer many services and it's used mostly for spinning up VMs. As such, running a cluster in DigitalOcean is very close to what we'd experience when running a cluster **on-prem**, be it **bare metal or VMs** created with, for example, VMWare.

You will be able to choose between **Ubuntu** and **CentOS** as operating systems. For storage, the instructions explain setup of a Kubernetes **NFS** client. We'll use Digital Ocean's load balancer. The logic behind its setup should be applicable to any other load balancer.

Throughout the document, we'll have sets of validations aimed at confirming that the cluster is set up correctly and can be used to install **CJE**. Feel free to jump straight into validations if you already have an operational cluster. The validations are focused on **RBAC** and **ServiceAccounts**, **load balancers** and **Ingress**, and **storage**.

Once we're confident that the cluster works as expected, we'll proceed with the CJE installation. We'll create CJOC, a managed master, and a job that will run in a separate Namespace.

We'll try to do as much work as possible through CLI. In a few cases, it will not be possible (or practical) to accomplish some tasks through a terminal window, so we'll have to jump into UIs. Hopefully, that will be only for very brief periods. The reason for insisting on CLI over UI, lies in the fact that commands are easier to reproduce and lead us towards automation. More importantly, I have a medical condition that results in severe pain when surrounded with many colors. The only medically allowed ones are black, white, and green. Unfortunately, most UIs are not designed for people with disabilities like mine.

At some later date, this document will be extended with the following items. Feel free to suggest additional ones.

- HAProxy as external LB
- nginx as external LB
- Ceph storage
- Gluster storage
- Basic CNI networking
- Flannel networking
- Calico networking
- Weave networking

Requirements

We'll need a few prerequisites first.

Please make sure that you have the following items.

- [kubectl](#): Used for communication with a Kubernetes cluster.
- [jq](#): Used for formatting and filtering of JSON outputs.
- [ssh-keygen](#): Used for generating SSH keys required for accessing nodes.
- [GitBash](#) (if Windows): Used for compatibility with other operating systems. Please use it **only if you are a Windows user**. It is available through [Git](#) setup.
- [DigitalOcean account](#): That's where we'll create a cluster.
- [doctl](#): CLI used for interaction with DigitalOcean API.

Creating A Rancher Server

The first step towards having a Kubernetes cluster is to have a Rancher server. Once it's up and running, we'll be able to use it to spin a Kubernetes cluster.

We'll need a DigitalOcean token that will allow us to authenticate with its API. Please open the API Tokens screen.

```
open "https://cloud.digitalocean.com/settings/api/tokens"
```

If you are a **Windows user**, you might not be able to use `open` command to interact with your browser. If that's the case, please replace `open` with `echo`, copy the output, and paste it into a new tab of your favorite browser.

Please type `cje` as the token name and click the *Generate Token* button. This is the first and the last time you will be able to see the token through DigitalOcean UI. Please store it somewhere safe. We'll need it soon.

Next, we'll create an SSH key that will allow us to enter into the virtual machines we'll create soon.

Please execute the command that follows.

```
ssh-keygen -t rsa
```

Please type `cje` as the file name. Feel free to answer to all the other question with the enter key.

Now that we have the SSH key, we can upload it to DigitalOcean. But, before we do that, we need to authenticate first.

```
doctl auth init
```

If this is the first time you're using `doctl`, you will be asked for the authentication token we created earlier.

The output should be similar to the one that follows.

```
Using token [...]
Validating token... OK
```

We can upload the SSH key with the `ssh-key create` command.

Please execute the command the follows.

```
doctl compute ssh-key create cje \
  --public-key "$(cat cje.pub)"
```

We created a new SSH key in DigitalOcean and named it `cje`. The content of the key was provided with the `--public-key` argument.

The output should be simialr to the one that follows.

ID	Name	FingerPrint
21418650	cje	28:f8:51:f0...

We'll need the ID of the new key. Instead of copying and pasting it from the output, we'll execute a query that will retrieve the ID from DigitalOcean. That way, we can retrieve it at any time instead of saving the output of the previous command.

```
KEY_ID=$(doctl compute ssh-key list \
```

```
| grep cje \  
| awk '{print $1}'\  
  
echo $KEY_ID
```

We executed `ssh-key list` command that retrieve all the SSH keys available in our DigitalOcean account. Further on, we used `grep` to filter the result so that only the key named `cje` is output. Finally, we used `awk` to output only the first column that contains the ID we're looking for.

The output of the latter command should be similar to the one that follows.

```
21418650
```

Next, we need to find out the ID of the image we'll use to create a VM that will host Rancher.

If your operating system of choice is **Ubuntu**, please execute the command that follows.

```
DISTRIBUTION=ubuntu-18-04-x64
```

Otherwise, if you prefer CentOS, the command is as follows.

```
DISTRIBUTION=centos-7-x64
```

Now matter which operating system we prefer, the important thing to note is that we have the environment variable `DISTRIBUTION` that holds the `slug` we can use to find out the ID of the image we'll use. Slug is DigitalOcean term that, in this context, describes the name of a distribution.

Now we can retrieve the ID of the image we'll use.

```
IMAGE_ID=$(doctl compute \  
  image list-distribution \  
  -o json \  
  | jq ".[] | select(.slug==\"$DISTRIBUTION\").id")  
  
echo $IMAGE_ID
```

The command retrieved the list of all the distributions and sent the output to `jq` which, in turn, filtered it so that only the ID of the image that matches our desired distribution is retrieved.

The output of the latter command should be similar to the one that follows.

```
34487567
```

Now we are finally ready to create a new VM that will soon host our Rancher server.

```
doctl compute droplet create rancher \  
  --enable-private-networking \  
  --image $IMAGE_ID \  
  --size s-2vcpu-4gb \  
  --region nyc3 \  
  --ssh-keys $KEY_ID
```

We executed a `compute droplet create` command that created a Droplet (DigitalOcean name for a node or a VM). We named it `rancher`, enabled private networking, and set the image to the ID we retrieved previously. We used `s-2vcpu-4gb` VM size that provides 2 CPUs and 4 GB RAM. The server will run in New York 3 region for no particular reason besides the fact that we had to choose one. Finally, we specify the SSH key ID we retrieved earlier so that we can enter into the newly created VM and complete the installation.

The output is as follows.

ID	Name	Public	IPv4	Private	IPv4	Public	IPv6	Memory	VCPU	Disk	Region	Image
96650503	rancher							4096	2	80	nyc1	Ubuntu

Please note that your ID will be different as we'll as the `Image` if you chose CentOS as your operating system of choice.

Next, we need to retrieve the IP of the new droplet (VM).

Please execute the command the follows.

```
RANCHER_IP=$(doctl compute droplet list \  
  -o json | \  
  jq -r '.[0].ipv4_address')
```

```
jq -r '.[0] | select(.name=="rancher").networks.v4[0].ip_address')  
echo $RANCHER_IP
```

We retrieved the list of all the droplets (VMs) in JSON format and sent the output to `jq`. It filtered the results so that only `rancher` is retrieved and output the IP address. We stored the final output as `RANCHER_IP` variable.

The output of the latter command will differ from one case to another. Mine is as follows.

```
208.68.39.72
```

Now we can enter into the droplet.

```
ssh -i cje root@$RANCHER_IP
```

Rancher runs as a container. So, our first step towards setting it up is to install Docker.

Please execute only the commands that match your operating system.

If you chose **Ubuntu**, the commands are as follows.

```
apt update  
apt install -y docker.io
```

If, on the other hand, your operating system is **CentOS**, the commands are as follows.

```
yum install -y \  
    yum-utils \  
    device-mapper-persistent-data \  
    lvm2  
  
yum-config-manager --add-repo \  
    https://download.docker.com/linux/centos/docker-ce.repo  
  
yum install -y docker-ce  
  
systemctl start docker
```

No matter the operating system of choice, we'll validate that Docker was installed correctly by outputting its version.

```
docker version
```

The output is as follows.

```
Client:
Version:      18.03.1-ce
API version:  1.37
Go version:   go1.9.5
Git commit:   9ee9f40
Built:        Thu Apr 26 07:20:16 2018
OS/Arch:      linux/amd64
Experimental: false
Orchestrator: swarm

Server:
Engine:
Version:      18.03.1-ce
API version:  1.37 (minimum version 1.12)
Go version:   go1.9.5
Git commit:   9ee9f40
Built:        Thu Apr 26 07:23:58 2018
OS/Arch:      linux/amd64
Experimental: false
```

Now we are ready to set up Rancher.

```
docker run -d \
  --restart=unless-stopped \
  -p 80:80 -p 443:443 \
  rancher/server:preview
```

We used `-d` to run the container in background (detached). We specified `restart` strategy to `unless-stopped` which will guarantee that Docker will make sure that the container is running even if the process inside it fails. Please note that this strategy does not make your Rancher fault tolerant. If the node hosting our Rancher server goes down, we'd loose everything. However, for the purpose of this exercise, a single Rancher container should be enough.

We published ports `80` and `443` , even though we do not have SL certificates.

Finally, we're using `preview` tag since, at the time of this writing, Rancher 2 is still not

production ready. Yet, version two brings a complete overhaul and it would be pointless to invest time in setting up Rancher 1.x.

Let's exit the node and test whether Rancher indeed works.

```
exit
```

We'll use [nip.io](#) to generate valid domain for Rancher, as well as for the CJE later on. The service provides a wildcard DNS for any IP address. It extracts IP from the nip.io subdomain and sends it back in the response. For example, if we generate 192.168.99.100.nip.io, it'll be resolved to 192.168.99.100. We can even add sub-sub domains like something.192.168.99.100.nip.io, and it would still be resolved to 192.168.99.100. It's a simple and awesome service that quickly became an indispensable part of my toolbox.

```
RANCHER_ADDR=$RANCHER_IP.nip.io  
echo $RANCHER_ADDR
```

The output of the latter command should be similar to the one that follows.

```
208.68.39.72.nip.io
```

Now we can finally open Rancher UI in browser.

```
open "https://$RANCHER_ADDR"
```

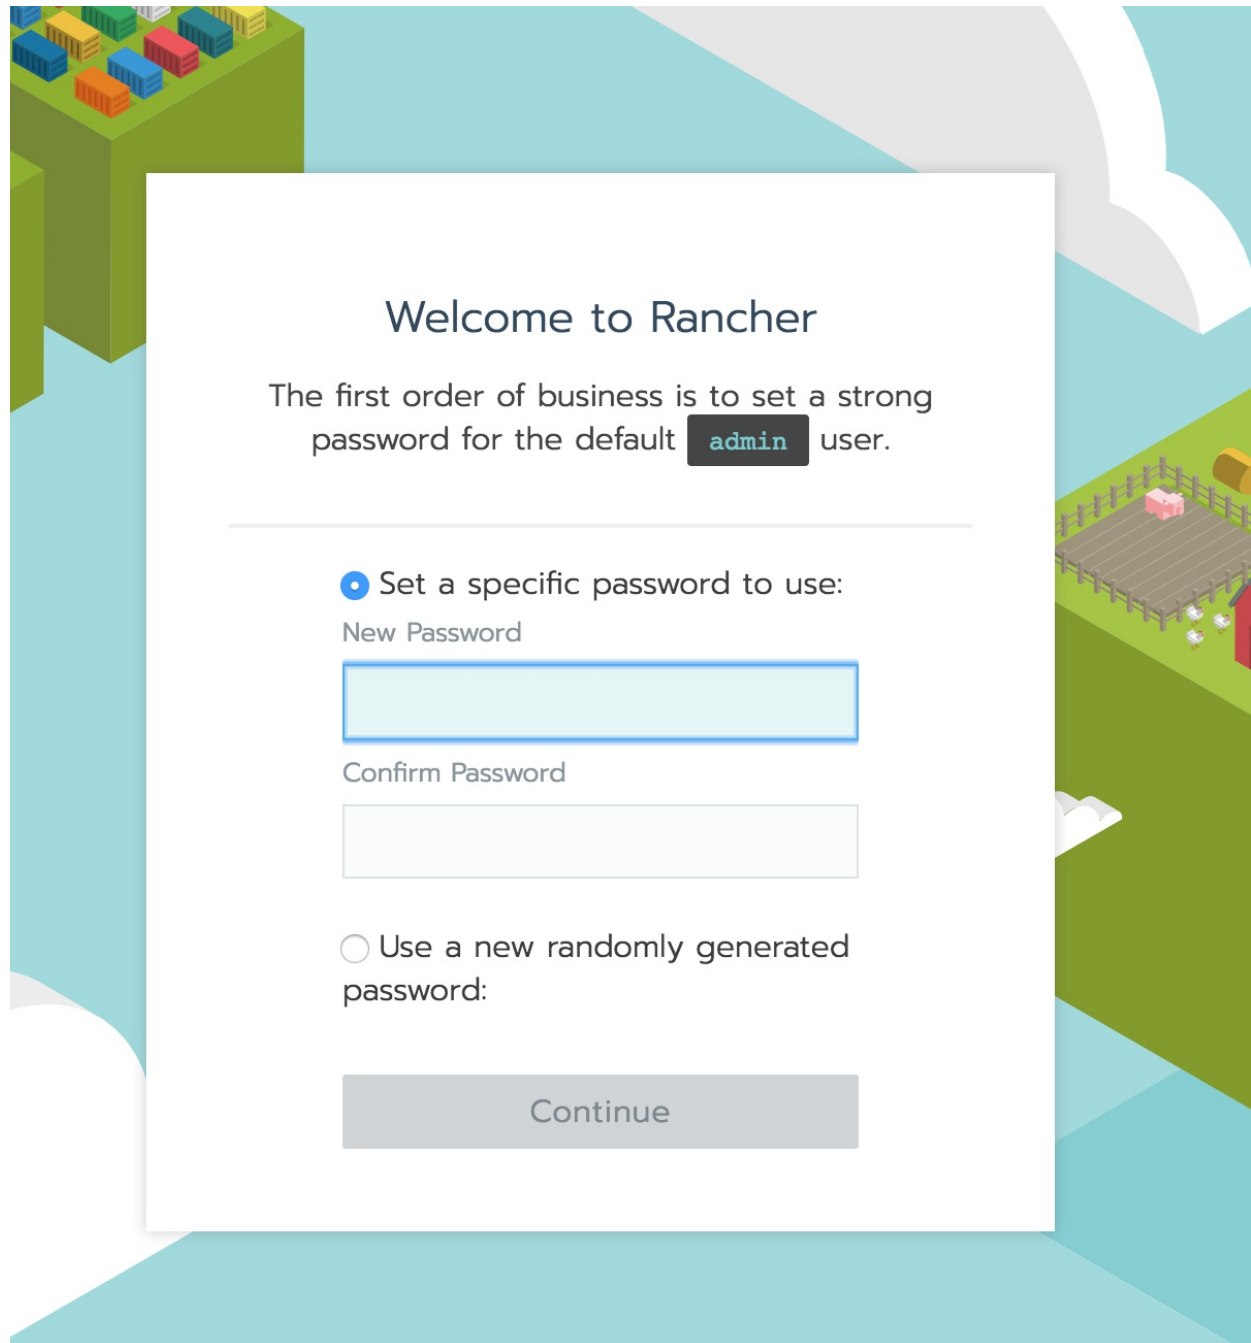
If you are a *Windows* user, Git Bash might not be able to use the `open` command. If that's the case, replace the `open` command with `echo`. As a result, you'll get the full address that should be opened directly in your browser of choice.

If you were quick, you might have seen the message stating that *this site can't be reached*. If that's the case, it means that Rancher process is still not up-and-running. Please wait a few moments and refresh the screen.

Since we do not have the certificates, you will probably see a warning or an error. If using Chrome, you'll have to click the *ADVANCED* link, followed with *Proceed to [YOUR_IP].nip.io (unsafe)*. If you prefer a different browser, the process should be similar and I'm sure you'll

know how to proceed.

You should see Rancher's welcome screen asking you to create the initial administrative password. Feel free to choose any, and type of twice. Once finished, please click the *Continue* button.

The image shows the Rancher welcome screen. It has a white central panel with a light blue background. The background of the entire screen features a stylized illustration of a farm with green hills, a blue sky, and various farm animals like cows and sheep. The text on the screen is as follows:

Welcome to Rancher

The first order of business is to set a strong password for the default `admin` user.

☒ Set a specific password to use:

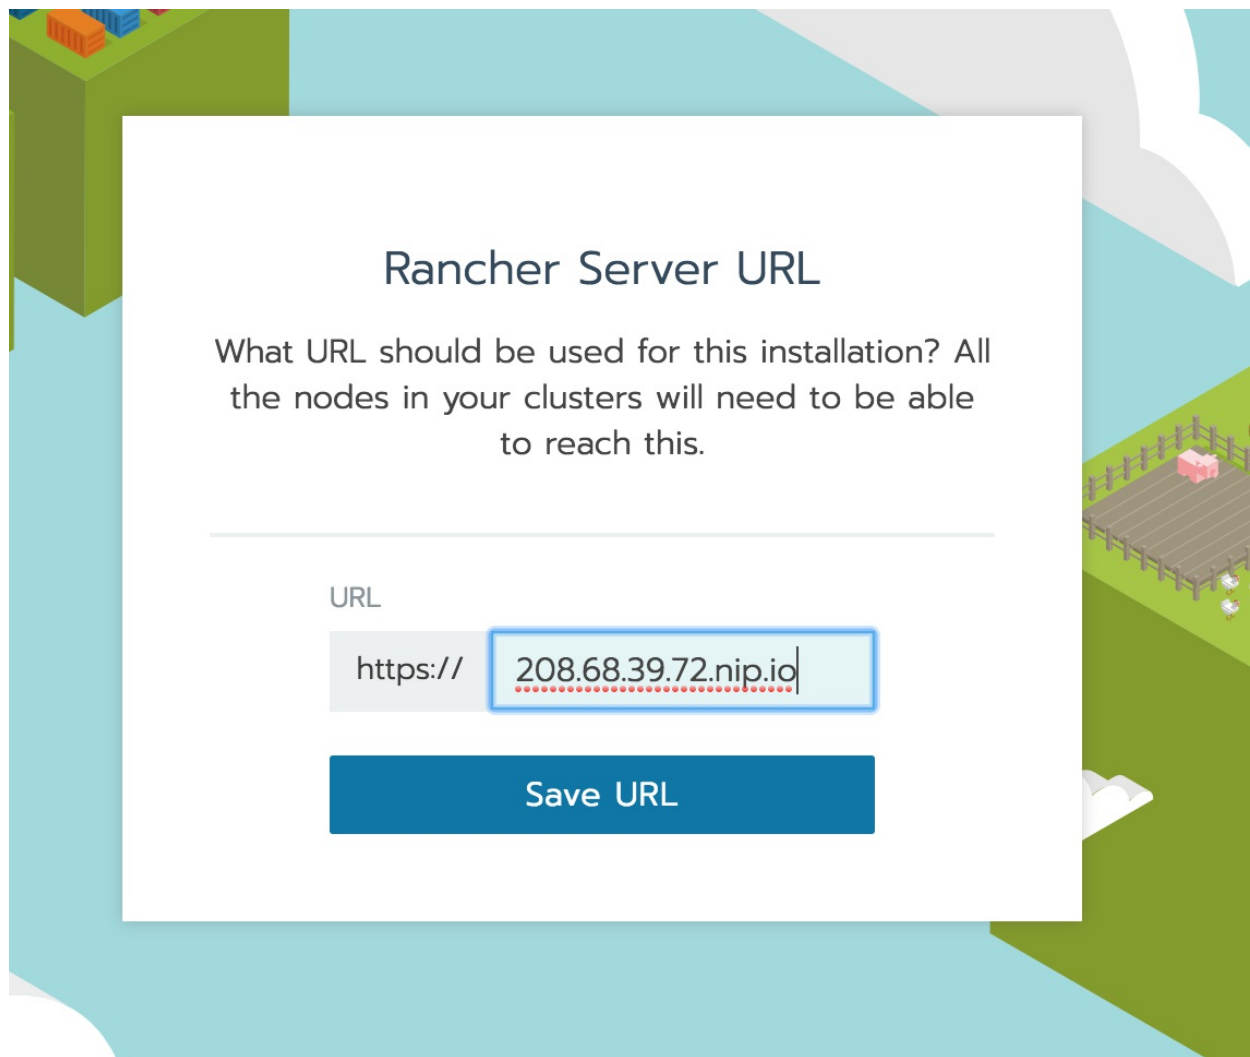
New Password

Confirm Password

☐ Use a new randomly generated password:

Continue

Next, you will be asked to specify *Rancher Server URL*. It is already pre-populated with the `nip.io` domain we used and all you have to do is click the *Save URL* button.



Creating An NFS Server

Before we proceed and create a cluster using Rancher, we'll need to set up an NFS server that will serve as storage for state generated by our applications.

Feel free to skip this whole section if you do not plan to use NFS.

For simplicity, we'll set up an NFS server on the same node where Rancher is running. Please don't do that for in "real-world" situations. Rather, you should have one or more separate NFS servers dedicated to a Kubernetes cluster.

The first order of business is to SSH into the node where we'll install and set up NFS.

```
NFS_SERVER_ADDR=$RANCHER_IP  
  
ssh -i cje root@$NFS_SERVER_ADDR
```

The steps required to install NFS server differ from one operating system to another. Please execute only the instructions matching your favorite OS.

If your operating system of choice is **Ubuntu**, please execute the commands that follow.

```
apt-get update  
  
apt-get install -y nfs-kernel-server  
  
NFS_USER=nobody:nogroup
```

On the other hand, if you prefer **CentOS**, the commands are as follows.

```
systemctl enable nfs-server.service  
  
systemctl start nfs-server.service  
  
NFS_USER=nfsnobody:nfsnobody
```

No matter the OS, the last command created the environment variable `NFS_USER` that defines the user and the group that will use soon to give a directory permissions that will allow any client to write files.

Please note that we are not setting up firewall and that we're assuming that any node can mount the NFS server we're about to set up.

Next, we'll create a directory we'll export and make the NFS user with wide permission own it.

```
mkdir /var/nfs/cje -p  
  
chown $NFS_USER /var/nfs/cje
```

We'll need to configure `/etc/exports` file that will be used by the NFS server to decide where to store the files, write mode, whether it should be synched, and so on.

```
echo "/var/nfs/cje *(rw,sync,no_subtree_check)" \  
| tee -a /etc/exports
```

The only thing left is to restart NFS server (if using Ubuntu), or to export NFS table (if using

CentOS).

If you're OS of choice is **Ubuntu**, please execute the command that follows.

```
systemctl restart nfs-kernel-server
```

CentOS users should export NFS table with the command that follows.

```
exportfs -a
```

Our NFS server is now up and running. Before we start using it, we'll need a Kubernetes cluster. For now, please exit the VM.

```
exit
```

Creating A Kubernetes Cluster

When using Rancher, the only option for creating a Kubernetes cluster is through its UI. So, we'll start by opening Rancher's home screen.

```
open "https://$RANCHER_ADDR"
```

Please click the *Add Cluster* button, and select *DigitalOcean*. Type *cje* as the *Cluster Name*.

Next, we will define the nodes that will constitute our cluster. Please scroll down to the *Node Pool* section.

First we'll create master nodes. We'll need three if we are to create a fault tolerant cluster. Master nodes need to have quorum before committing any decision. If we'd have two master nodes, and one of them fails, we'd lose quorum since over 50% is required.

Type *master* as the *Name Prefix*, set *Count* to *3*, and select *etcd* and *Control* check boxes.

The only thing left, when defining master nodes, is to decide which node specifications we'll use. We do that through Node Templates.

Please click the *Add Node Template* button (icon with a + sign). You'll have to type or, more likely, paste the DigitalOcean token we created earlier. Once inside the template screen, please select *New York 3* as the *Region*, select *2 GB RAM*, *60 GB Disk*, *2 vCPUs* as *Droplet Size*, and type *k8s* as the *Name*. The purpose of each of those fields should be self-explanatory.

The only thing left to do, before we create the new template, is to choose the image. If you prefer **Ubuntu**, please select **Ubuntu 16.04 x64**. Similarly, **CentOS** users should pick **CentOS 7.X x64**.

Once finished, click the *Create* button.

Next, we'll repeat a similar process to define worker nodes. Those are the nodes where our applications will run.

Please click the *Add Node Pool* button. A new row will be added. Type *worker* as the *Name Prefix*, set *Count* to 3, and make sure that the *Worker* checkbox is selected.

We are almost done. All that's left is to click the *Create* button, and be patient.

Name Prefix	Count		Template		etcd	Control	Worker	
master	1	x	k8s	+	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	-
worker	3	x	k8s	+	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	-

Number of nodes required:

☒ 1, 3, or 5

☒ 1 or more


☒ 1 or more

+ Add Node Pool


Create

Cancel

After instructing Rancher to create a new cluster, we're redirected to the *Clusters* screen. At the moment, *cje* should be the only cluster managed by Rancher and its state should be *Provisioning*. Please click the *cje* link in the *Cluster Name* column.



Global
Clusters
Node Drivers
Catalogs
Users
Settings
Security




Clusters Add Cluster

Delete
Search


<input type="checkbox"/>	State	Cluster Name	Provider	Nodes	CPU	RAM	
<input type="checkbox"/>	Provisioning	cje	DigitalOcean	4	n/a	n/a	⋮

Waiting for etcd and controlplane nodes to be registered

Now we need to be patient and wait until everything is green. It should take around 10 minutes.



cje
Cluster
Nodes
Storage
Projects/Namespaces
Members
Tools



Dashboard: cje Launch kubectl Kubeconfig File

Provider DigitalOcean
Version v1.10.1
Nodes 4

16%
CPU
1.3 of 8 used

2%
Memory
0.1 of 6.8 GiB used

4%
Pods
16 of 440 used

etcd

Controller Manager

Scheduler

Nodes

Once our new cluster is fully operational, we can configure out local `kubectl`. We'll create a separate config file only for this cluster.

```
vim do-kube-config.yaml
```

Feel free to change `vim` to your favorite editor. All we'll need is to paste some content, so any editor you're comfortable with should do.

Please go back to Rancher UI, click the *Kubeconfig File* button, followed with a click to the *Copy to Clipboard* link.

```
nk0nBrq2hQxTfYdubQmZlnZ0LwZvMv1hXVSKVAZ0hHnZx0x0mES0fNlZLTP0bHMIenBMZkKCTZQX
1JxT2oxRXF1NE82V01VaVZsMnE00VZPaWREUXZvZ1haNUhkZHJBaz1zNke9PQotLS0tLUVORCBDR\
VJUSUZJQ0FURS0tLS0tCg==\"

users:
- name: \"user-wtqct\"
  user:
    token: \"kubeconfig-user-wtqct:2mbnjfvfcvtq6r7t5vxxm6mzbd9fvt9tfvnkrcw2wdhhv1vxfs1kq\"

contexts:
- name: \"cje\"
  context:
    user: \"user-wtqct\"
    cluster: \"cje\"

current-context: \"cje\"
```

 Copy to Clipboard

Then [download](#) (if needed) and run `kubect1`

Close

Next, we need to go back to the terminal window where we have `do-kube-config.yaml` opened in our favorite editor, paste the content we just copied, and save the file. If you're using `vim`, you should press `i` to enter into the insert mode, paste the copied content, press `esc` button to exit the insert mode, press colon (`:`) to enter into the command mode, and type `wq` to save and quit. If you prefer something other than `vim`, I'm sure you'll know how to do the same operations (paster, save, and exit).

Kubect1 uses environment variable `KUBECONFIG` to know the location of the config file, so we'll need to export it with the path to the file we just created.

```
export KUBECONFIG=$PWD/do-kube-config.yaml
```

For now, we'll do only a rudimentary validation of the cluster by retrieving the nodes. If everything worked as expected, we should have six of them, three dedicated to master nodes (`controlplane` and `etcd`), and the other three as workers.

```
kubect1 get nodes
```

The output is as follows.

NAME	STATUS	ROLES	AGE	VERSION
master1	Ready	controlplane,etcd	1h	v1.10.1
master2	Ready	controlplane,etcd	1h	v1.10.1
master3	Ready	controlplane,etcd	1h	v1.10.1
worker1	Ready	worker	1h	v1.10.1
worker2	Ready	worker	1h	v1.10.1
worker3	Ready	worker	1h	v1.10.1

We're almost finished with creating the cluster. All that's left is to create two namespaces. We'll use `cjoc` to host CloudBees Operations Center as well as for the managed masters we'll spin up later. The second namespace will be called `build`. That where our builds will run Pods. With such separation, we'll try to mitigate some of the risks builds might introduce to CJOC and masters.

```
kubectl create ns cjoc

kubectl create ns build
```

Validating RBAC

In this section, we'll validate whether RBAC is configured correctly. Specifically, we'll check whether Pods are denied the right to communicate with KubeAPI with proper service accounts.

We'll start by creating two namespaces.

```
kubectl create ns test1

kubectl create ns test2
```

Assuming that RBAC is properly set up (we'll test that soon), we'd need to create a few service accounts that will allow processes running inside Pods in those namespaces, communicate with Kube API. As a test, we'll define roles that will allow us to do almost any Pod-related operation from within another Pod. At the same time, we'll need to confirm that we are not allowed to create Pods (or any other Kubernetes resource) in any other Namespace.

Let's take a quick look at a definition that creates a service account, a few roles, and a few role bindings.


```
curl https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sa/pods-all.yml
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: pods-all
  namespace: test1

---

kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: pods-all
  namespace: test1
rules:
- apiGroups: [""]
  resources: ["pods", "pods/exec", "pods/log"]
  verbs: ["*"]

---

kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: pods-all
  namespace: test2
rules:
- apiGroups: [""]
  resources: ["pods", "pods/exec", "pods/log"]
  verbs: ["*"]

---

apiVersion: rbac.authorization.k8s.io/v1beta1
kind: RoleBinding
metadata:
  name: pods-all
  namespace: test1
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: pods-all
subjects:
- kind: ServiceAccount
  name: pods-all

---

apiVersion: rbac.authorization.k8s.io/v1beta1
kind: RoleBinding
metadata:
  name: pods-all
```

```
namespace: test2
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: pods-all
subjects:
- kind: ServiceAccount
  name: pods-all
  namespace: test1
```

The ServiceAccount `pods-all` will be created in the Namespace `test1`. As a result, only the Pods in the same Namespace will be able to use it.

Further on, we have two Roles named `pods-all`. Their specs are the same except for the fact that one is in the Namespace `test1`, and the other in `test2`. The permissions behind those roles provide the ability to do almost any operation on Pods, and no other resource type.

Finally, the last two resources are RoleBindings. Each references the Role `pods-all` in their respective Namespaces. The major difference is that the second RoleBinding is in the Namespace `test2` but it relates to the ServiceAccount `pods-all` in the Namespace `test1`.

As a result of that definition, we should be able to create Pods in the Namespace `test1` and from within the containers in those Pods, create new Pods, output their logs, and so on. Any Pod operation within the two Namespaces initiated through a process in a Pod running inside the Namespace `test1` should be allowed. RBAC works in a way that we specify only what is allowed, and not what isn't. So, any other operation should not be allowed.

Let's apply the resources from the `pods-all.yml` file.

```
kubectl apply \
  -f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sa/pods-all.yml
```

The output is as follows.

```
serviceaccount "pods-all" created
role "pods-all" created
role "pods-all" created
rolebinding "pods-all" created
rolebinding "pods-all" created
```

Next, we'll create a Pod that will allow us to test the assumptions. Let's take a look at the

definition first.

```
curl https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sa/kubect1-test2.yaml
```

The output is as follows.

```
apiVersion: v1
kind: Pod
metadata:
  name: kubect1
  namespace: test1
spec:
  serviceAccountName: pods-all
  containers:
  - name: kubect1
    image: vfarcic/kubect1
    command: ["sleep"]
    args: ["100000"]
```

It defines a Pod with a container based on `vfarcic/kubect1` image. It contains `kubect1` that will allow us (or not) to spin up additional Pods. The Pod uses the ServiceAccount `pods-all` and will run in the namespace `test1`.

Let's apply the definition.

```
kubect1 apply \
  -f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sa/kubect1-test2.
```

Please wait a few moments until the image is pulled and the Pod is up-and-running.

Next, we'll enter into the container that forms the Pod.

```
kubect1 -n test1 exec -it kubect1 -- sh
```

Can we create a Pod in the `test2` Namespace? Let's see.

```
kubect1 -n test2 \
```

```
run new-test \  
--image=alpine \  
--restart=Never \  
sleep 10000
```

The output is as follows.

```
pod "new-test" created
```

Judging by the output, the Pod was created and, therefore, our ServiceAccount has the sufficient permissions.

Can we list the Pods in the `test2` Namespace?

```
kubectl -n test2 get pods
```

The output is as follows.


NAME	READY	STATUS	RESTARTS	AGE
new-test	1/1	Running	0	17s

Checking only the operations we should be allowed to execute gives us only part of the picture. Those commands do not prove that we are NOT allowed to issue requests to Kube API outside of those specified in the Roles. To test that, we'll try to create a Pod outside Namespaces `test1` and `test2`.

```
kubectl -n default get pods
```

The output is as follows.

```
Error from server (Forbidden): pods is forbidden: User "system:serviceaccount:test1
```



That confirms that the ServiceAccount `pods-all` cannot list pods in the namespace `"default"`.

We're done with a quick RBAC validation. We'll exit the `kubectl` Pod and delete the test Namespaces.

```
exit

kubectl delete ns test1 test2
```

Creating A Load Balancer

TODO: Continue

DigitalOcean LB

```
WORKER_IDS=$(doctl compute \
  droplet list -o json | \
  jq -r '.[ ] | select(.name | startswith("worker")).id' \
  | tr '\n' ',' | tr -d ' ')

echo $WORKER_IDS
```

```
104.131.98.145,104.236.49.97,104.131.106.27,
```

```
WORKER_IDS=${WORKER_IDS: :-1}

echo $WORKER_IDS
```

```
104.131.98.145,104.236.49.97,104.131.106.27
```

```
kubectl -n ingress-nginx get svc default-http-backend
```

```
doctl compute load-balancer create \
  --droplet-ids $WORKER_IDS \
  --forwarding-rules "entry_protocol:tcp,entry_port:80,target_protocol:tcp,target_
  --health-check protocol:http,port:80,path:/healthz,check_interval_seconds:10,re
  --name cje \
  --region nyc3
```

```
LB_IP=$(doctl compute load-balancer \
  list -o json | jq -r '.[0].ip')
```

```
echo $LB_IP
```

```
45.55.123.83
```

```
LB_ADDR=$LB_IP.nip.io
```

Validating Load Balancer And Ingress

```
kubectl apply \
  -f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/ingress/go-demo-2
```

```
ingress "go-demo-2" created
deployment "go-demo-2-db" created
service "go-demo-2-db" created
deployment "go-demo-2-api" created
service "go-demo-2-api" created
```

```
kubectl rollout status \
  deploy go-demo-2-api
```

```
Waiting for rollout to finish: 0 of 3 updated replicas are available...
Waiting for rollout to finish: 1 of 3 updated replicas are available...
Waiting for rollout to finish: 2 of 3 updated replicas are available...
deployment "go-demo-2-api" successfully rolled out
```

```
curl -i "http://$LB_ADDR/demo/hello"
```

```
HTTP/1.1 200 OK
Server: nginx/1.13.8
Date: Thu, 07 Jun 2018 00:30:51 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 14
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains;

hello, world!
```

```
curl -i "http://$LB_ADDR/this/does/not/exist"
```

```
HTTP/1.1 404 Not Found
Server: nginx/1.13.8
Date: Thu, 07 Jun 2018 00:31:23 GMT
Content-Type: text/plain; charset=utf-8
Content-Length: 21
Connection: keep-alive
Strict-Transport-Security: max-age=15724800; includeSubDomains;

default backend - 404
```

```
kubectl delete \
  -f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/ingress/go-demo-2
```

```
ingress "go-demo-2" deleted
deployment "go-demo-2-db" deleted
service "go-demo-2-db" deleted
deployment "go-demo-2-api" deleted
service "go-demo-2-api" deleted
```

Creating StorageClasses

NFS

```
kubectl -n cjoc create \
  -f https://raw.githubusercontent.com/kubernetes-incubator/external-storage/mast
```

```
serviceaccount "nfs-client-provisioner" created
```

```
kubectl create \
  -f https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master
```

```
clusterrole "nfs-client-provisioner-runner" created
```

```
curl https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master
  | sed -e "s@namespace: default@namespace: cjoc@g" \
  | kubectl create -f -
```

```
clusterrolebinding "run-nfs-client-provisioner" created
```

```
curl https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master
  | sed -e "s@10.10.10.60@$NFS_SERVER_ADDR@g" \
  | sed -e "s@/ifs/kubernetes@/var/nfs/cje@g" \
  | sed -e "s@fuseim.pri/ifs@cloudbees.com/cje-nfs@g" \
  | kubectl -n cjoc create -f -
```

```
deployment "nfs-client-provisioner" created
```

```
kubectl -n cjoc \
  rollout status \
  deploy nfs-client-provisioner
```

```
Waiting for rollout to finish: 0 of 1 updated replicas are available...
```



```
deployment "nfs-client-provisioner" successfully rolled out
```

```
curl https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master  
| sed -e "s@managed-nfs-storage@cje-storage@g" \  
| sed -e "s@fuseim.pri/ifs@cloudbees.com/cje-nfs@g" \  
| kubectl -n cjoc create -f -
```

```
storageclass "cje-storage" created
```

```
curl https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master  
| sed -e "s@managed-nfs-storage@cje-storage@g" \  
| kubectl -n cjoc create -f -
```

```
persistentvolumeclaim "test-claim" created
```

```
kubectl create \  
-n cjoc \  
-f https://raw.githubusercontent.com/kubernetes-incubator/external-storage/mast
```

```
pod "test-pod" created
```

```
ssh -i cje root@$NFS_SERVER_ADDR \  
"ls -l /var/nfs/cje"
```

```
drwxrwxrwx. 2 nfsnobody nfsnobody 21 Jun  7 00:44 cjoc-test-claim-pvc-d5fa9106-69eb
```

```
ssh -i cje root@$NFS_SERVER_ADDR \  
  "ls -l /var/nfs/cje/cjoc-test*"
```

```
-rw-r--r--. 1 nfsnobody nfsnobody 0 Jun 7 00:44 SUCCESS
```

```
kubectl delete \  
  -n cjoc \  
  -f https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master
```

```
pod "test-pod" deleted
```

```
kubectl delete \  
  -n cjoc \  
  -f https://raw.githubusercontent.com/kubernetes-incubator/external-storage/master
```

```
persistentvolumeclaim "test-claim" deleted
```

```
ssh -i cje root@$NFS_SERVER_ADDR \  
  "ls -l /var/nfs/cje"
```

```
drwxrwxrwx. 2 nfsnobody nfsnobody 21 Jun 7 00:44 archived-cjoc-test-claim-pvc-d5fa
```

Validating StorageClasses

```
kubectl -n cjoc apply \  
  -f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sts/cje-test.yml
```

```
statefulset "test" created
service "test" created
```

```
kubectl -n cjoc exec test-0 \
  -- touch /tmp/something

kubectl -n cjoc exec test-0 \
  -- ls -l /tmp
```

```
-rw-r--r--    1 nobody   nobody           0 Jun  7 01:04 something
```

```
kubectl -n cjoc delete pod test-0
```

```
pod "test-0" deleted
```

```
kubectl -n cjoc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nfs-client-provisioner-69688c76dd-b2bjj	1/1	Running	0	33m
test-0	1/1	Terminating	0	3m

```
kubectl -n cjoc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nfs-client-provisioner-69688c76dd-b2bjj	1/1	Running	0	34m
test-0	1/1	Running	0	9s

```
kubectl -n cjoc exec test-0 \
```

```
-- ls -l /tmp
```

```
-rw-r--r--    1 nobody   nobody           0 Jun  7 01:04 something
```

```
kubectl -n cjoc delete \  
-f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sts/cje-test.yml
```

```
statefulset "test" deleted  
service "test" deleted
```

```
kubectl -n cjoc \  
delete pvc test-data-test-0
```

```
persistentvolumeclaim "test-data-test-0" deleted
```

```
kubectl -n build apply \  
-f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sts/cje-test.yml
```

```
statefulset "test" created  
service "test" created
```

```
kubectl -n build exec test-0 \  
-- touch /tmp/something  
  
kubectl -n build exec test-0 \  
-- ls -l /tmp
```

```
-rw-r--r--    1 nobody   nobody           0 Jun  7 01:11 something
```

```
kubectl -n build delete pod test-0
```

```
pod "test-0" deleted
```

```
kubectl -n build get pods
```

NAME	READY	STATUS	RESTARTS	AGE
test-0	1/1	Terminating	0	58s

```
kubectl -n build get pods
```

NAME	READY	STATUS	RESTARTS	AGE
test-0	1/1	Running	0	3s

```
kubectl -n build exec test-0 \  
-- ls -l /tmp
```

```
-rw-r--r--    1 nobody   nobody           0 Jun  7 01:11 something
```

```
kubectl -n build delete \  
-f https://raw.githubusercontent.com/vfarcic/k8s-specs/master/sts/cje-test.yml
```

```
statefulset "test" deleted  
service "test" deleted
```

```
kubectl -n build \  

```

```
delete pvc test-data-test-0
```

```
persistentvolumeclaim "test-data-test-0" deleted
```

TODO: Speed tests

Installing CJE

```
open "https://downloads.cloudbees.com/cje2/latest/"
```

```
# Copy the link address of cje2 Kubernetes release
```

```
RELEASE_URL=[...]
```

```
curl -o cje.tgz $RELEASE_URL
```

```
tar -xvf cje.tgz
```

```
x cje2_2.121.1.2_kubernetes/  
x cje2_2.121.1.2_kubernetes/cjoc-external-masters.yml  
x cje2_2.121.1.2_kubernetes/INSTALLATION.md  
x cje2_2.121.1.2_kubernetes/SCALING.md  
x cje2_2.121.1.2_kubernetes/TROUBLESHOOTING.md  
x cje2_2.121.1.2_kubernetes/PROXY.md  
x cje2_2.121.1.2_kubernetes/ANALYTICS.md  
x cje2_2.121.1.2_kubernetes/cje.yml
```

```
cd cje2_*
```

```
ls -l
```

```
-rw-r--r--  1 vfarcic  staff   158  6 Jun 07:06 ANALYTICS.md  
-rw-r--r--  1 vfarcic  staff  5144  6 Jun 07:06 INSTALLATION.md  
-rw-r--r--  1 vfarcic  staff  1498  6 Jun 07:06 PROXY.md  
-rw-r--r--  1 vfarcic  staff  1466  6 Jun 07:06 SCALING.md  
-rw-r--r--  1 vfarcic  staff  4950  6 Jun 07:06 TROUBLESHOOTING.md  
-rw-r--r--  1 vfarcic  staff 11884  6 Jun 07:06 cje.yml  
-rw-r--r--  1 vfarcic  staff   509  6 Jun 07:06 cjoc-external-masters.yml
```

```
cat cje.yml

kubectl create ns jenkins

cat cje.yml \
| sed -e \
"s@https://cje.example.com@http://cje.example.com@g" \
| sed -e \
"s@cje.example.com@$LB_ADDR@g" \
| sed -e \
"s@ssl-redirect: \"true\"@ssl-redirect: \"false\"@g" \
| sed -e \
"s@# storageClassName: some-storage-class@storageClassName: cje-storage@g" \
| kubectl --namespace cjoc \
apply -f -
```

```
serviceaccount "cjoc" created
role "master-management" created
rolebinding "cjoc" created
configmap "cjoc-config" created
configmap "cjoc-configure-jenkins-groovy" created
statefulset "cjoc" created
service "cjoc" created
ingress "cjoc" created
serviceaccount "jenkins" created
role "pods-all" created
rolebinding "jenkins" created
configmap "jenkins-agent" created
```

```
kubectl -n cjoc \
  rollout status sts cjoc
```

```
Waiting for 1 pods to be ready...
statefulset rolling update complete 1 pods at revision cjoc-578b8fd6b4...
```

```
kubectl -n cjoc get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/nfs-client-provisioner	1	1	1	1	51m

NAME	DESIRED	CURRENT	READY	AGE
------	---------	---------	-------	-----

```
rs/nfs-client-provisioner-69688c76dd 1 1 1 51m
```

NAME	DESIRED	CURRENT	AGE
statefulsets/cjoc	1	1	1m

NAME	READY	STATUS	RESTARTS	AGE
po/cjoc-0	1/1	Running	0	1m
po/nfs-client-provisioner-69688c76dd-b2bjj	1/1	Running	0	51m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/cjoc	ClusterIP	10.43.128.241	<none>	80/TCP,50000/TCP	1m

```
open "http://$LB_ADDR/cjoc"
```

```
kubectl -n cjoc exec cjoc-0 -- \
    cat /var/jenkins_home/secrets/initialAdminPassword
```

```
c9b802a409a04d1f8755e9d394749a58
```

```
# Copy the output and paste it into Jenkins UI field *Administrative password*
```

```
# Click the *Continue* button
```

```
# Click the *Use a license key* button
```

```
# Type your *License Key* and *License Certificate*
```

```
# Click the *OK* button
```

```
# Click the *Install suggested plugins* button
```


```
# Fill in the field in the *Create First Admin User* screen
```

```
# Click the *Save and Continue* button
```

```
# Click the *Start using Operations Center* button
```

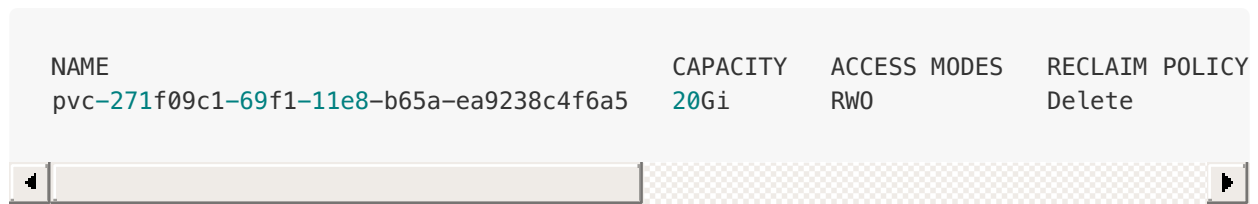
```
kubectl -n cjoc get pvc
```

NAME	STATUS	VOLUME	CAPACITY
jenkins-home-cjoc-0	Bound	pvc-271f09c1-69f1-11e8-b65a-ea9238c4f6a5	20Gi




```
kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY
pvc-271f09c1-69f1-11e8-b65a-ea9238c4f6a5	20Gi	RWO	Delete



Creating Managed Masters

TODO: Create

TODO: Fail and confirm that the state is preserved

Managed Agents (PodTemplates)

TODO

External Masters

TODO

External Agents

TODO

Validating CJE

Destroying The Cluster

TODO